

Lab 01: Data Cleaning

CS3300 Data Science

Learning Outcomes

1. Identify, access, load, and prepare (clean) a data set for a given problem.
2. Select, apply, and interpret appropriate visual and statistical methods to analyze distributions of individual variables.
3. Clean and transform data for analysis.
4. Communicate findings through generated data visualizations and reports.

Overview

In this lab, you are going to inspect and clean a data set of real estate transactions from California. You should prepare your results as a Jupyter notebook. In addition to code and plots, you should have text offering interpretations and explanations. Your notebook should be organized into sections with appropriate headers. The notebook and its code should be clean and polished.

Instructions

1. Loading the Data and Initial Assessment

- a. Load the `Sacramentorealestatetransactions.csv` file as a `DataFrame`.
- b. Using the output of the initial `head()` and `info()` commands, in a dedicated markdown cell in your notebook - describe the data. What are the variables? What are their inferred types? Do any of the columns have null values?

2. Representing Categorical Variables

- a. Sometimes a variable can either be represented as an integer or categorical variable. Count the number of unique values for the streets, zip codes, and beds. Do you think it is more appropriate to represent these three variables as categorical or integer variables? Why or why not?
- b. Convert the following variables to categorical variables: city, state, zip, beds, baths, type. You may find these two links from Pandas useful: [astype](#) and [basic types](#).

3. Exploring Continuous Variables

- a. Plot histograms of the square footage, latitudes, and longitudes.

b. Do you notice “odd” patterns in any of the plots?

4. Exploring Categorical Variables

a. Plot the beds, baths, type, state, city, and zip codes as count (bar) plots. ([seaborn countplot](#) might be helpful).

b. Is there anything “odd” about the numbers of beds and baths?

5. Engineering New Variables – Part I

Entries with 0 square footage are empty lots. Encoding these cases with values of 0 lead to two different interpretations of the square footage variable. This is a good candidate for creating a new boolean variable.

a. Create a new boolean variable called "empty_lot". This variable should have a value of true if the square footage of a record is 0. Otherwise, it should have a value of false.

b. Create a count (bar) plot for the empty_lot variable.

6. Engineering New Variables – Part II

a. Count of the number of unique values for the addresses (streets variable). Do you think this variable is useful for analysis or as a feature for a ML model in its current form?

b. Street types (e.g., avenue, street, way) can indicate whether a road will be quiet or busy, is in a commercial or suburban area, etc. The street types can be extracted from the address. Using the head() command, look at the first 20 records to identify address patterns. You can check [here](#) for a list of possible of street suffix abbreviations.

c. Write a function get_street_type(address) that will return the street type (as a String) of an address.

d. Use that function to create a new categorical variable of street types. Count the unique elements and print them out. If it looks like any addresses were parsed incorrectly (e.g., a number is returned instead of a street type), update your function in part c. Repeat until you parse everything correctly.

e. Plot the street types as a count (bar) plot.

7. Identifying Potential Label Variables

This data set can be used for both regression and classification problems.

- c. Identify a variable which would make a good label (output) variable for a regression problem.
- d. Identify a variable which would make a good label (output) variable for a classification problem.

8. Identify and remove an outlier record from the dataset.

There is categorical variable that has some entry as “unknown”. Identify which records or rows these are and remove them from your data frame.

9. Save the Cleaned Data Set

Save your cleaned data set as a CSV file using `to_csv` with the `index=False` option.

Submission Instructions

Save the Jupyter notebook as a PDF and upload that file through Canvas.

Rubric

Followed submission instructions	5%
Formatting: Report is polished and clean. No unnecessary code. Section headers are used. Plots are described and interpreted using text. The report contains an introduction and conclusion.	5%
Loading and Describing Data	5%
Representing Categorical Variables	10%
Cleaning Continuous Variables	10%
Cleaning Categorical Variables	10%
Engineering New Variables – Part I	10%
Engineering New Variables – Part II (including correct parsing of all street types)	20%
Identification of Dependent Variables	10%
Reasonable effort is made to answers all questions (e.g., 3-4 sentences each)	10%
Technical/Organizational/Presentation Proficiency	5%